


# Computer Networks

## Lecture 6a: Interconnecting LANs

Based on slides from D. Choffnes Northeastern U. and P. Gill from StonyBrook University  
Revised Autumn 2015 by S. Laki

# Just Above the Data Link Layer

2




- Bridging
  - How do we connect LANs?
- Function:
  - Route packets between LANs
- Key challenges:
  - Plug-and-play, self configuration
  - How to resolve loops

# Recap

3


- Originally, Ethernet was a broadcast technology



- Pros: Simplicity
  - Hardware is stupid and cheap
- Cons: No scalability
  - More hosts = more collisions = pandemonium

# Bridging the LANs

4



- Bridging limits the size of collision domains
  - Vastly improves scalability
  - Question: could the whole Internet be one bridging domain?
- Tradeoff: bridges are more complex than hubs
  - Physical layer device vs. data link layer device
  - Need memory buffers, packet processing hardware, routing tables


# Bridges

5

- ❑ Original form of Ethernet switch

❑ Consistent with IEEE 802 LAN Local Area Network

- ❑
  1. Forwarding of frames
  2. Learning of (MAC) Addresses
  3. Spanning Tree Algorithm (to handle loops)
- No hardware or software changes on hosts/hubs
- Should not impact existing LAN operations



# Frame Forwarding Tables

6

- Each bridge maintains a forwarding table

MAC Address	Port	Age
00:00:00:00:00:AA	1	1 minute
00:00:00:00:00:BB	2	7 minutes
00:00:00:00:00:CC	3	2 seconds
00:00:00:00:00:DD	1	3 minutes



# Learning Addresses

7


- ❑ Manual configuration is possible, but...
  - Time consuming
  - Error Prone
  - Not adaptable (hosts may get added or removed)

- ❑ Instead, learn addresses using heuristic

Delete old entries after a timeout

- Look at the **source** of frames that arrive on each port


MAC Address	Port	Age
00:00:00:00:00:AA	1	0 minutes
00:00:00:00:00:BB	2	0 minutes



# The Danger of Loops

8

- ❑ <Src=AA, Dest=DD>
- ❑ This continues to infinity
  - How do we stop this?
- ❑ Remove loops from the topology
  - Without physically unplugging cables
- ❑ 802.1 (LAN architecture) uses an algorithm to build and maintain a **spanning tree** for routing






# Spanning Tree Definition

9

- A subset of edges in a graph that:
  - Span all nodes
  - Do not create any cycles
- This structure is a tree



# 802.1 Spanning Tree Approach

10

1. Elect a bridge to be the root of the tree
  2. Every bridge finds shortest path to the root
  3. Union of these paths becomes the spanning tree
- 
- Bridges exchange Configuration Bridge Protocol Data Units (BPDUs) to build the tree
    - Used to elect the root bridge
    - Calculate shortest paths
    - Locate the next hop closest to the root, and its port
    - Select ports to be included in the spanning trees

# Determining the Root

11

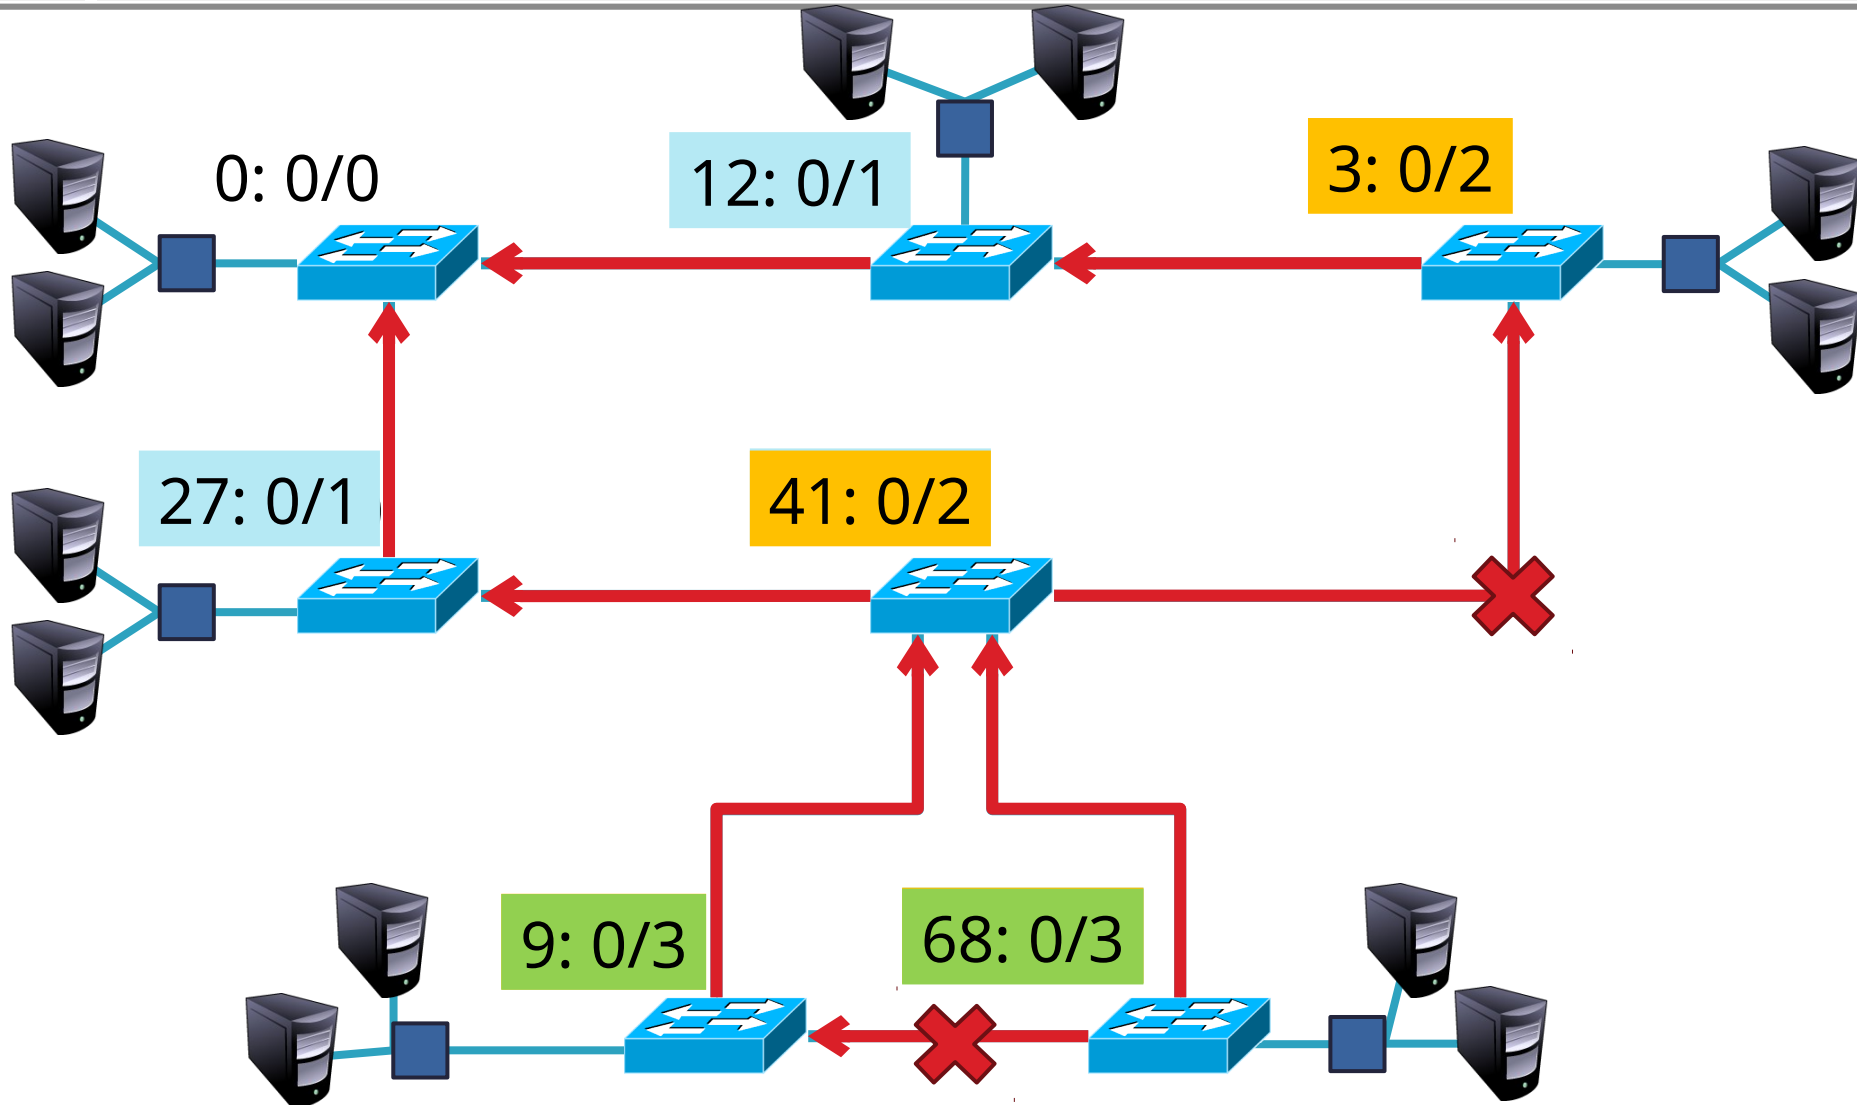
- Initially, all hosts assume they are the root
- Bridges broadcast BPDUs:

Bridge ID	Root ID	Path Cost to Root
-----------	---------	-------------------

- Based on received BPDUs, each switch chooses:
  - A new root (smallest known Root ID)
  - A new root port (what interface goes towards the root)
  - A new designated bridge (who is the next hop to root)

# Spanning Tree Construction

12



# Bridges vs. Switches

13

- ❑ Bridges make it possible to increase LAN capacity
  - Reduces the amount of broadcast packets
  - No loops
- ❑ Switch is a special case of a bridge
  - Each port is connected to a **single** host
    - Either a client machine
    - Or another switch
  - Links are full duplex
  - Simplified hardware: no need for CSMA/CD!
  - Can have different speeds on each port

# Switching the Internet

14

- Capabilities of switches:
  - Network-wide routing based on MAC addresses
  - Learn routes to new hosts automatically
  - Resolve loops
- Could the whole Internet be one switching domain?

NO

# Limitations of MAC Routing

15

- ❑ Inefficient
  - Flooding packets to locate unknown hosts
- ❑ Poor Performance
  - Spanning tree does not balance load
  - Hot spots
- ❑ Extremely Poor Scalability
  - Every switch needs every MAC address on the Internet in its routing table!
- ❑ IP addresses these problems (next ...)

# Computer Networks


## Lecture 6b: Network Layer

Based on slides from D. Choffnes Northeastern U. and P. Gill from StonyBrook University  
Revised Autumn 2015 by S. Laki



# Network Layer

17




- Function:
  - Route packets end-to-end on a network, through multiple hops
- Key challenge:
  - How to represent addresses
  - How to route packets
    - Scalability
    - Convergence

# Routers, Revisited

18

- How to connect multiple LANs?
- LANs may be incompatible
  - Ethernet, Wifi, etc...
- Connected networks form an **internetwork**
  - The Internet is the best known example



## Internet Service Model

- Best-effort (i.e. things may break)
- Store-and-forward datagram network

Lowest common denominator

work)

### □ Service Model

- What gets sent?
- How fast will it go?
- What happens if there are failures?
- Must deal with **heterogeneity**
  - Remember, every network is different

- ❑ Addressing
  - ❑ Class-based
  - ❑ CIDR
- ❑ IPv4 Protocol Details
  - ❑ Packed Header
  - ❑ Fragmentation
- ❑ IPv6

# Possible Addressing Schemes

21

## □ Flat

- e.g. each host is identified by a 48-bit MAC address
- Router needs an entry for every host in the world
  - Too big
  - Too hard to maintain (hosts come and go all the time)
  - Too slow (more later)

## □ Hierarchy


- Addresses broken down into segments
- Each segment has a different level of specificity

# Example: Telephone Numbers

22

1-617-373- 3278

Very General



Northeastern University

West Village G  
Room 254

Updates are Local

Very specific

# IP Addressing

24


- IPv4: 32-bit addresses
  - Usually written in dotted notation, e.g. 192.168.21.76
  - Each number is a byte

	0	8	16	24	31
Decimal	192	168	21	76	
Hex	C0	A8	15	4C	
Binary	11000000	10101000	00010101	01001100	

# IP Addressing and Forwarding

25


- Routing Table Requirements
  - For every possible IP, give the next hop
  - But for 32-bit addresses,  $2^{32}$  possibilities!
- Hierarchical address scheme
  - Separate the address into a network and a host





# Classes of IP Addresses

26



# How Do You Get IPs?

27

- IP address ranges controlled by IANA

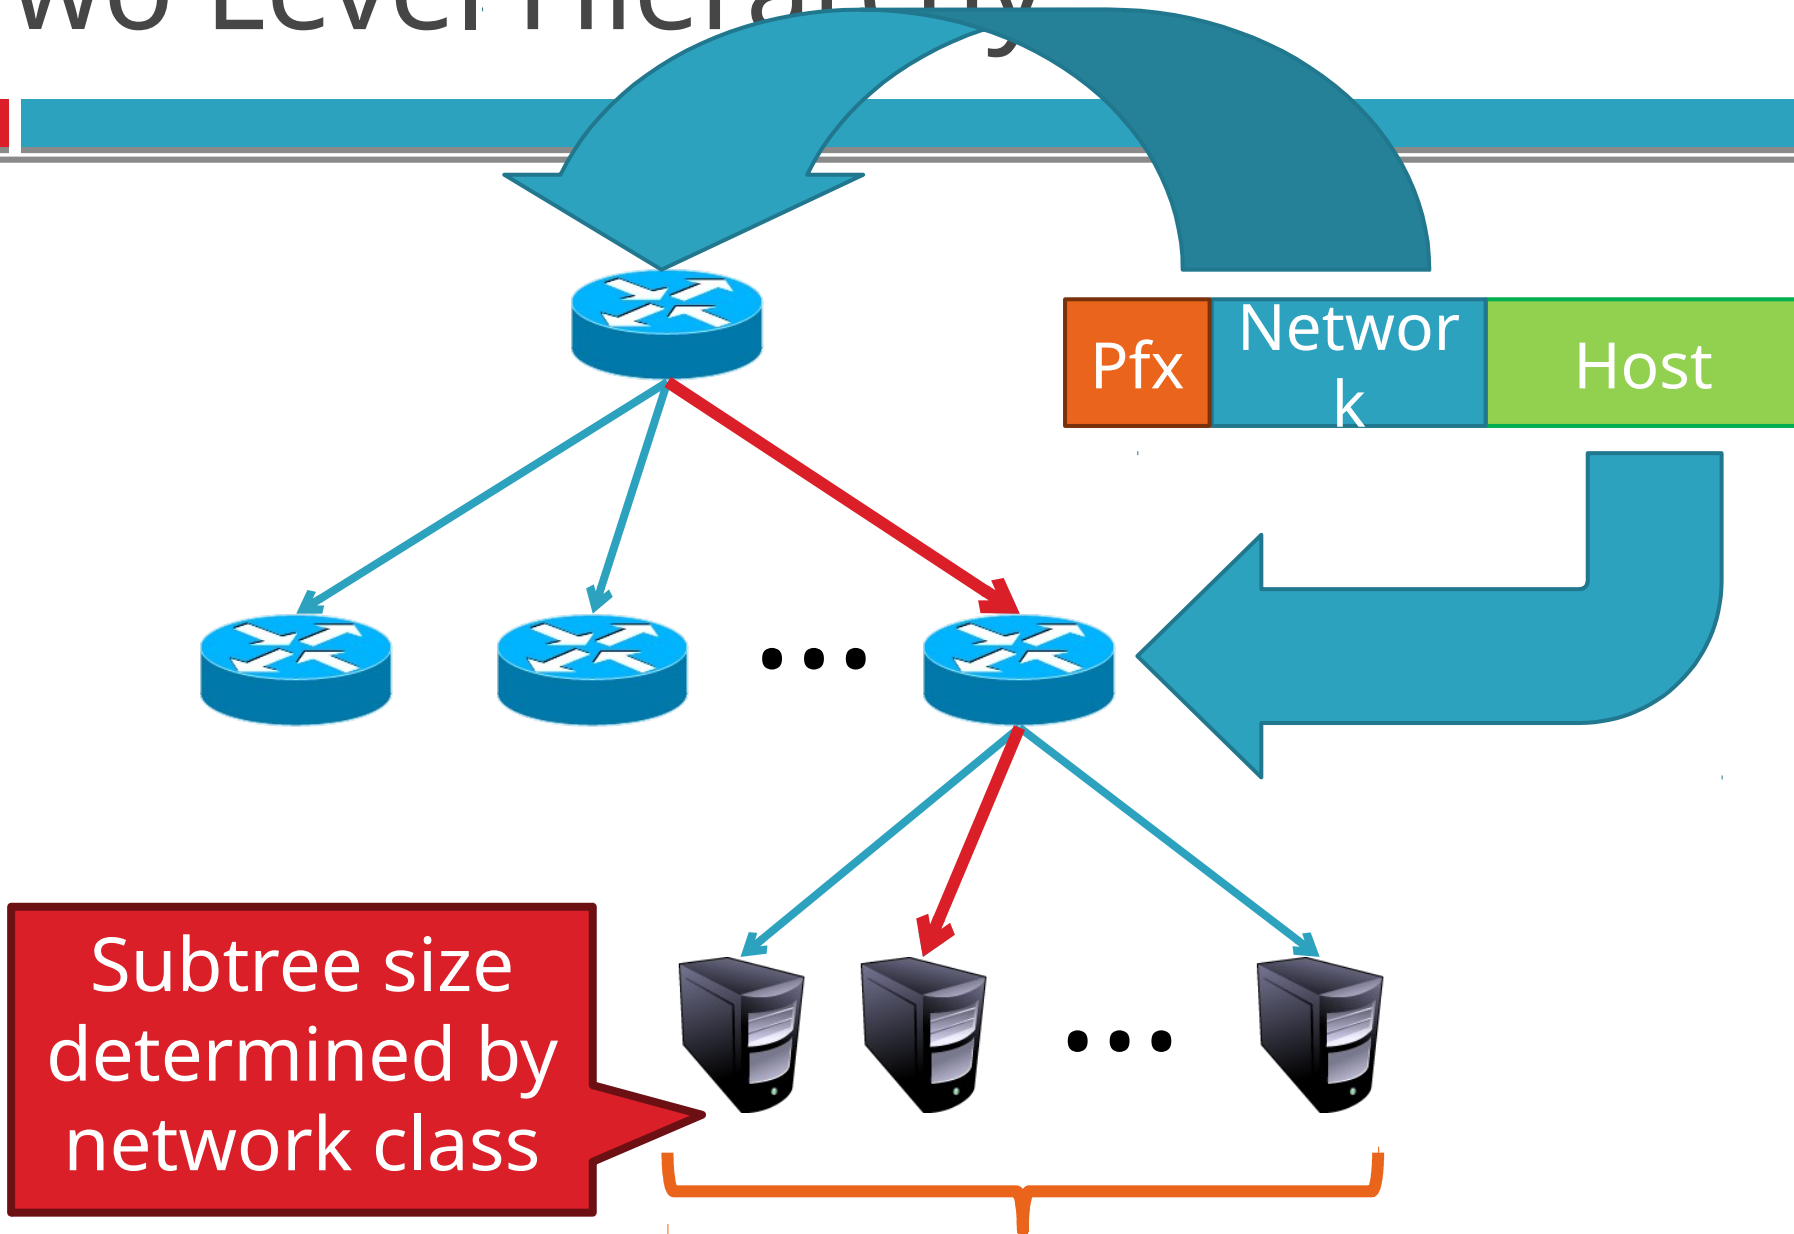


Internet Assigned Numbers Authority

- Internet Assigned Number Authority
  - Roots go back to 1972, ARPANET, UCLA
  - Today, part of ICANN
- IANA grants IPs to regional authorities
    - ARIN (American Registry of Internet Numbers) may grant you a range of IPs
    - You may then advertise routes to your new IP range
    - There are now secondary markets, auctions, ...

# Two Level Hierarchy

28



# Class Sizes

29

Way too big

Class	Prefix Bits	Network Bits	Number of Classes	Hosts per Class
A	1	7	$2^7 - 2 = 126$ (0 and 127 are reserved)	$2^{24} - 2 = 16,777,214$ (All 0 and all 1 are reserved)
B	2	14	$2^{14} - 2 = 16,382$ (0 and 1 are reserved)	$2^{16} - 2 = 65,534$ (All 0 and all 1 are reserved)
C	3	21	$2^{21} = 2,097,152$ (0 and 1 are reserved)	$2^8 - 2 = 254$ (All 0 and all 1 are reserved)
D	4	28	$2^{28} = 268,435,456$ (0 and 1 are reserved)	$2^0 = 1$ (All 0 and all 1 are reserved)
E	5	31	$2^{31} = 2,147,483,648$ (0 and 1 are reserved)	$2^{-1} = 0.5$ (All 0 and all 1 are reserved)

Too many network IDs

Too small to be useful

# Subnets

30

- Problem: need to break up large A and B classes
- Solution: add another layer to the hierarchy
  - From the outside, appears to be a single network
    - Only 1 entry in routing tables
  - Internally, manage multiple subnetworks
    - Split the address range using a **subnet mask**



Subnet Mask: 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0

# Subnet Example

31

## □ Extract network:

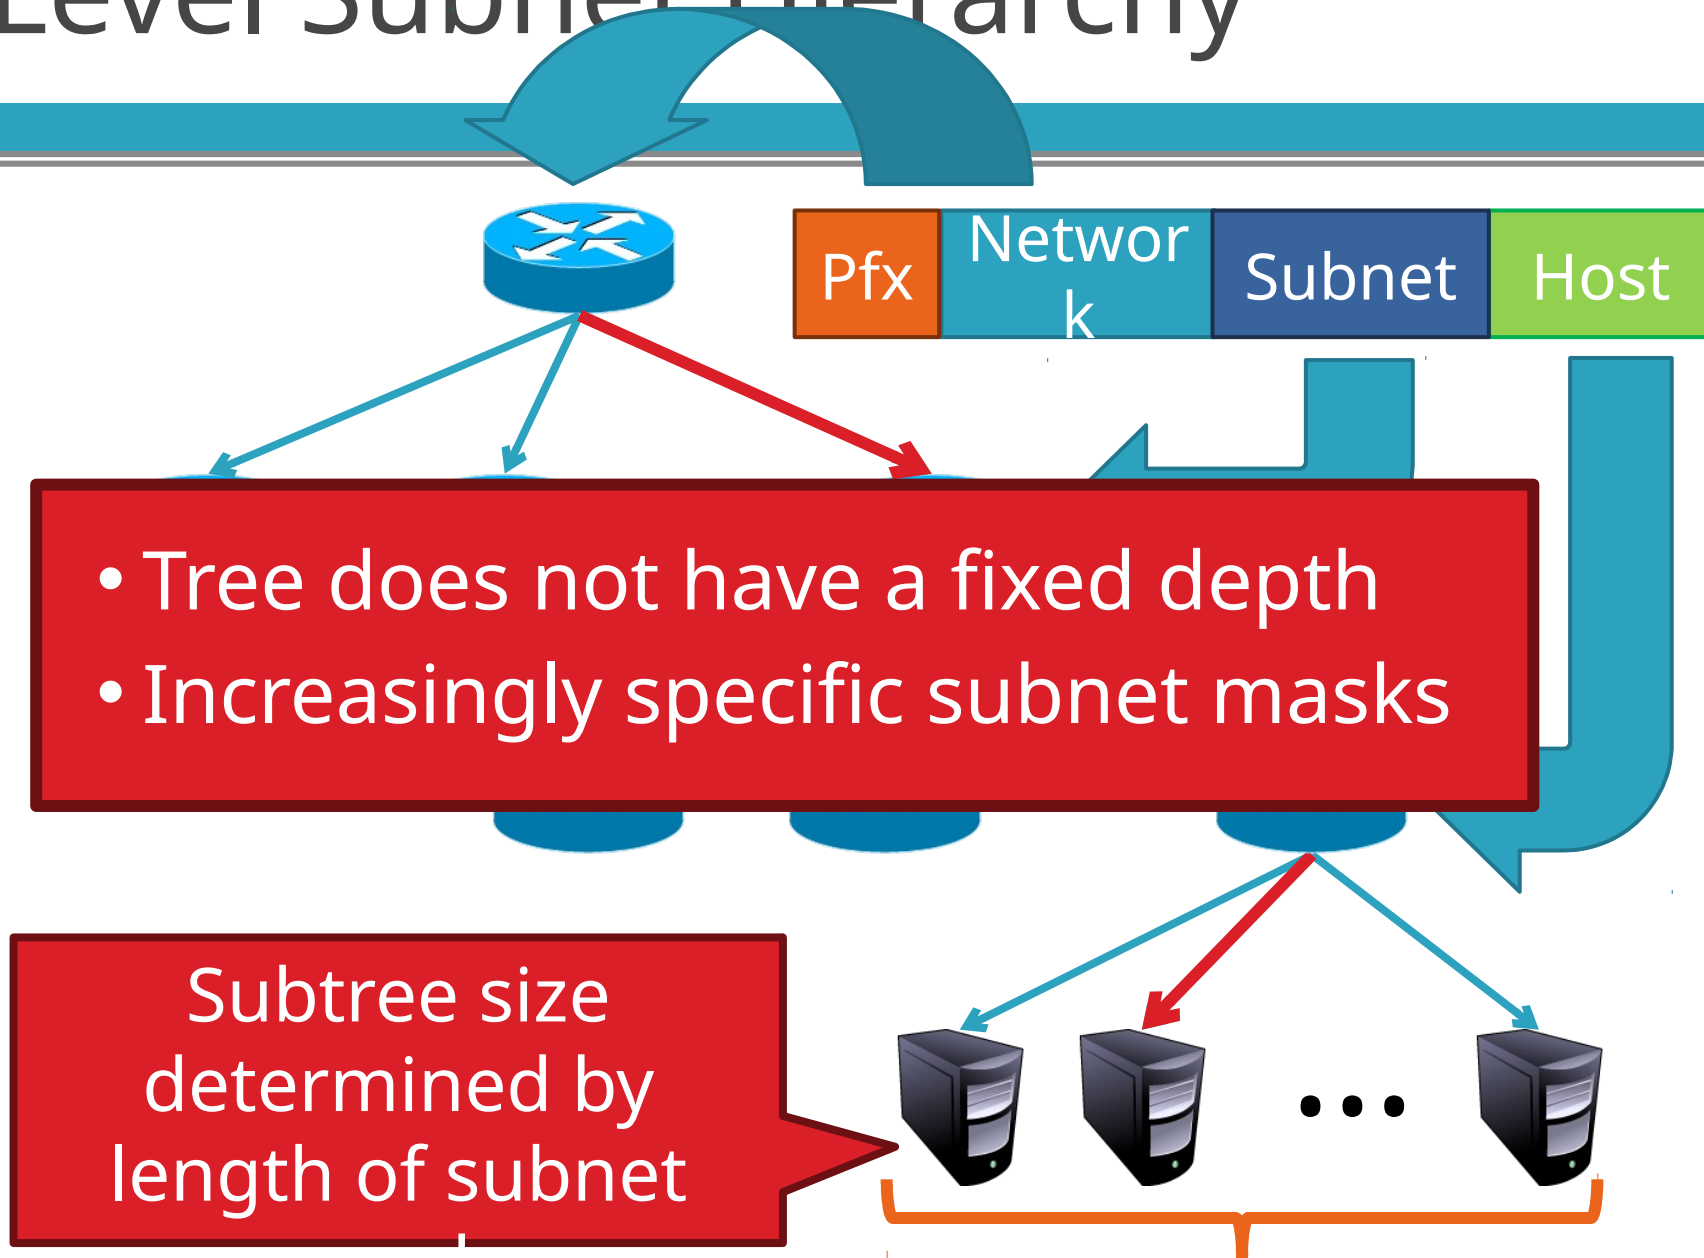
IP Address:	10110101 11011101 01010100 01110010
Subnet Mask:	& 11111111 11111111 11000000 00000000
Result:	10110101 11011101 01000000 00000000

## □ Extract host:

IP Address:	10110101 11011101 01010100 01110010
Subnet Mask:	& ~(11111111 11111111 11000000 00000000)
Result:	00000000 00000000 00010100 01110010

# N-Level Subnet Hierarchy

32



# Example Routing Table

33

Address Pattern	Subnet Mask	Destination Router
0.0.0.0	0.0.0.0	Router 4
18.0.0.0	255.0.0.0	Router 2
128.42.0.0	255.255.0.0	Router 3
128.42.128.0	255.255.128.0	Router 5
128.42.222.0	255.255.255.0	Router 1

❑ Question: 128.42.222.198 matches all four rows

- Which router do we forward to?

❑ Longest prefix matching

- Use the row with the longest number of 1's in the mask
- This is the **most specific match**



# Subnetting Revisited

34

- Question: does subnetting solve all the problems of class-based routing?

NO

- Classes are still too coarse
  - Class A can be subnetted, but only 126 available
  - Class C is too small
  - Class B is nice, but there are only 16,398 available
- Routing tables are still too big
  - 2.1 million entries per router

# Classless Inter Domain Routing


35

- CIDR, pronounced 'cider'
- Key ideas:
  - Get rid of IP classes
  - Use bitmasks for all levels of routing
  - **Aggregation** to minimize FIB (forwarding information base)
- Arbitrary split between network and host
  - Specified as a bitmask or prefix length
  - Example: Stony Brook
    - 130.245.0.0 with **netmask** 255.255.0.0
    - 130.245.0.0 / 16

# Aggregation with CIDR


36

- ❑ Original use: aggregating class C ranges
- ❑ One organization given contiguous class C ranges
  - Example: Microsoft, 207.46.192.\* – 207.46.255.\*
  - Represents  $2^6 = 64$  class C ranges
  - Specified as CIDR address 207.46.192.0/18



# Size of CIDR Routing Tables

37



- From [www.cidr-report.org](http://www.cidr-report.org)
- CIDR has kept IP routing table sizes in check
  - Currently ~500,000 entries for a complete IP routing table
  - Only required by backbone routers

# We had a special day in summer 2014!

38

- 512K day – August 12, 2014
- Default threshold size for IPv4 route data in older Cisco routers □ 512K routes
  - Some routers failed over to slower memory
    - RAM vs. CAM (content addressable memory)
  - Some routes dropped
- Cisco issues update in May anticipating this issue
  - Reallocated some IPv6 space for IPv4 routes
- Part of the cause
  - Growth in emerging markets
- <http://cacm.acm.org/news/178293-internet-routing-failures-bring-architecture-changes-back-to-the-table/fulltext>

# Takeaways

39

- ❑ Hierarchical addressing is critical for scalability
  - Not all routers need all information
  - Limited number of routers need to know about changes
- ❑ Non-uniform hierarchy useful for heterogeneous networks
  - Class-based addressing is too coarse
  - CIDR improves scalability and granularity
- ❑ Implementation challenges
  - Longest prefix matching is more difficult than schemes with no ambiguity