

# Computer Networks

## **Lecture 11-12: Transport layer Part II**

### **DNS, HTTP**

Based on slides from D. Choffnes Northeastern U. and P. Gill from StonyBrook University  
Revised Autumn 2015 by S. Laki

# Transport Layer

2



## □ Function:

- ▣ Demultiplexing of data streams

## □ Optional functions:

- ▣ Creating long lived connections
- ▣ Reliable, in-order packet delivery
- ▣ Error detection
- ▣ Flow and congestion control

## □ Key challenges:


- ▣ Detecting and responding to congestion
- ▣ Balancing fairness against high utilization

# What is Congestion?

3

- Load on the network is higher than capacity
  - ▣ Capacity is not uniform across networks
    - Modem vs. Cellular vs. Cable vs. Fiber Optics
- ▣ There are multiple flows competing for bandwidth

▣ Load




# Cong. Control vs. Cong. Avoidance

4

Congestion Avoidance:  
Stay left of the knee

Congestion Control:  
Stay left of the cliff



Congestion  
Collapse

# TCP Congestion Control

5

- Each TCP connection has a window
  - ▣ Controls the number of unACKed packets
- Sending rate is  $\sim \text{window}/\text{RTT}$
- Idea: vary the window size to control the send rate
- Introduce a congestion window at the sender
  - ▣ Congestion control is sender-side problem

# Two Basic Components

6

## 1. Detect congestion

- Packet dropping is most reliable signal
  - Delay-based methods are hard and risky
- How do you detect packet drops? ACKs
  - Timeout after not receiving an ACK
  - Several duplicate ACKs in a row (ignore for now)

## 2. Rate adjustment algorithm

- Modify *cwnd*
- Probe for bandwidth
- Responding to congestion

# Rate Adjustment

7

- Recall: TCP is ACK clocked
  - ▣ Congestion = delay = long wait between ACKs
  - ▣ No congestion = low delay = ACKs arrive quickly
- Basic algorithm
  - ▣ Upon receipt of ACK: increase *cwnd*
    - Data was delivered, perhaps we can send faster
    - *cwnd* growth is proportional to RTT
  - ▣ On loss: decrease *cwnd*
    - Data is being lost, there must be congestion
- Question: increase/decrease functions to use? !!!!

# Implementing Congestion Control

8


- ❑ Maintains three variables:
  - ❑ *cwnd*: congestion window
  - ❑ *adv\_wnd*: receiver advertised window
  - ❑ *ssthresh*: threshold size (used to update *cwnd*)
- ❑ For sending, use:  $wnd = \min(cwnd, adv\_wnd)$
- ❑ Two phases of congestion control
  1. Slow start ( $cwnd < ssthresh$ )
    - Probe for bottleneck bandwidth
  2. Congestion avoidance ( $cwnd \geq ssthresh$ )
    - AIMD



# Slow Start

9


- ❑ Goal: reach knee quickly
- ❑ Upon starting (or restarting) a connection
  - ❑  $cwnd = 1$
  - ❑  $ssthresh = adv\_wnd$
  - ❑ Each time a segment is ACKed,  $cwnd++$
- ❑ Continues until...
  - ❑  $ssthresh$  is reached
  - ❑ Or a packet is lost
- ❑ Slow Start is not actually slow
  - ❑  $cwnd$  increases exponentially



# Slow Start Example

10

- $cwnd$  grows rapidly
- Slows down when...
  - ▣  $cwnd \geq ssthresh$
  - ▣ Or a packet drops




# Congestion Avoidance

11

- Additive Increase Multiplicative Decrease (AIMD) mode
- *ssthresh* is lower-bound guess about location of the knee
- **If** *cwnd*  $\geq$  *ssthresh* **then**
  - each time a segment is ACKed
  - increment *cwnd* by  $1/cwnd$  ( $cwnd += 1/cwnd$ ).
- So *cwnd* is increased by one only if all segments have been acknowledged

# Congestion Avoidance Example

12



$cwnd = 1$


$cwnd = 2$

$cwnd = 4$

$ssthresh = 8$

$cwnd = 8$


$cwnd = 9$



# The Big Picture – TCP Tahoe

(the original TCP)

13



- ❑ UDP
- ❑ TCP
- ❑ Congestion Control
- ❑ **Evolution of TCP**
- ❑ Problems with TCP

# The Evolution of TCP


15

- ❑ Thus far, we have discussed TCP Tahoe
  - ▣ Original version of TCP
- ❑ However, TCP was invented in 1974!
  - ▣ Today, there are many variants of TCP
- ❑ Early, popular variant: TCP Reno
  - ▣ Tahoe features, plus...
  - ▣ Fast retransmit
    - 3 duplicate ACKs?  $\rightarrow$  retransmit (don't wait for RTO)
  - ▣ Fast recovery
    - On loss:  $\text{set } \text{cwnd} = \text{cwnd}/2$  ( $\text{ssthresh} = \text{new cwnd value}$ )

# TCP Reno: Fast Retransmit

16

- Problem: in Tahoe, if segment is lost, there is a long wait until the RTO
- Reno: retransmit after 3 duplicate ACKs






# TCP Reno: Fast Recovery

17

- After a fast-retransmit set  $cwnd$  to  $cwnd/2$ 
  - ▣ Also reset  $ssthresh$  to the new halved  $cwnd$  value
  - ▣ i.e. don't reset  $cwnd$  to 1
  - ▣ Avoid unnecessary return to slow start
  - ▣ Prevents expensive timeouts
- But when RTO expires still do  $cwnd = 1$ 
  - ▣ Return to slow start, same as Tahoe
  - ▣ Indicates packets aren't being delivered at all
  - ▣ i.e. congestion must be really bad

# Fast Retransmit and Fast Recovery

18



- ❑ At steady state,  $cwnd$  oscillates around the optimal window size
- ❑ TCP always forces packet drops

# Many TCP Variants...

19

- ❑ Tahoe: the original
  - ❑ Slow start with AIMD
  - ❑ Dynamic RTO based on RTT estimate
- ❑ Reno:
  - ❑ fast retransmit (3 dupACKs)
  - ❑ fast recovery ( $cwnd = cwnd/2$  on loss)
- ❑ NewReno: improved fast retransmit
  - ❑ Each duplicate ACK triggers a retransmission
  - ❑ Problem:  $>3$  out-of-order packets causes pathological retransmissions
- ❑ Vegas: delay-based congestion avoidance
- ❑ And many, many, many more...

# TCP in the Real World

20

- What are the most popular variants today?
  - ▣ Key problem: TCP performs poorly on high bandwidth-delay product networks (like the modern Internet)
  - ▣ Compound TCP (Windows)
    - Based on Reno
    - Uses two congestion windows: delay based and loss based
    - Thus, it uses a *compound* congestion controller
  - ▣ TCP CUBIC (Linux)
    - Enhancement of BIC (Binary Increase Congestion Control)
    - Window size controlled by cubic function
    - Parameterized by the time  $T$  since the last dropped packet

# High Bandwidth-Delay Product

21

- Key Problem: TCP performs poorly when
  - ▣ The capacity of the network (bandwidth) is large
  - ▣ The delay (RTT) of the network is large
  - ▣ Or, when bandwidth \* delay is large
    - $b * d = \text{maximum amount of in-flight data in the network}$
    - a.k.a. the bandwidth-delay product
- Why does TCP perform poorly?
  - ▣ Slow start and additive increase are slow to converge
  - ▣ TCP is ACK clocked
    - i.e. TCP can only react as quickly as ACKs are received
    - Large RTT  $\rightarrow$  ACKs are delayed  $\rightarrow$  TCP is slow to react

# Goals

22

- ❑ Fast window growth
  - ▣ Slow start and additive increase are too slow when bandwidth is large
  - ▣ Want to converge more quickly
- ❑ Maintain fairness with other TCP variants
  - ▣ Window growth cannot be too aggressive
- ❑ Improve RTT fairness
  - ▣ TCP Tahoe/Reno flows are not fair when RTTs vary widely
- ❑ Simple implementation


# Compound TCP Implementation

23

- Default TCP implementation in Windows
- Key idea: split *cwnd* into two separate windows
  - ▣ Traditional, loss-based window
  - ▣ New, delay-based window
- $wnd = \min(cwnd + dwnd, adv\_wnd)$ 
  - ▣ *cwnd* is controlled by AIMD
  - ▣ *dwnd* is the delay window
- Rules for adjusting *dwnd*:
  - ▣ If RTT is increasing, decrease *dwnd* ( $dwnd \geq 0$ )
  - ▣ If RTT is decreasing, increase *dwnd*
  - ▣ Increase/decrease are proportional to the rate of change

# Compound TCP Example

24



- Aggressiveness corresponds to changes in RTT
- Advantages: fast ramp up, more fair to flows with different RTTs
- Disadvantage: must estimate RTT, which is very challenging



# TCP CUBIC Implementation

25

- Default TCP implementation in Linux
- Replace AIMD with cubic function


$$W_{cubic} = C(T - K)^3 + W_{max} \quad (1)$$

C is a scaling constant, and  $K = \sqrt[3]{\frac{W_{max}\beta}{C}}$

- $\beta \rightarrow$  a constant fraction for multiplicative increase
- $T \rightarrow$  time since last packet drop
- $W_{max} \rightarrow$  cwnd when last packet dropped


# TCP CUBIC Implementation

26



# TCP CUBIC Example

27



- ❑ Less wasted bandwidth due to fast ramp up
- ❑ Stable region and slow acceleration help maintain fairness
  - ▣ Fast ramp up is more aggressive than additive increase
  - ▣ To be fair to Tahoe/Reno, CUBIC needs to be less aggressive

- ❑ UDP
- ❑ TCP
- ❑ Congestion Control
- ❑ Evolution of TCP
- ❑ Problems with TCP

# Issues with TCP

29

- The vast majority of Internet traffic is TCP
- However, many issues with the protocol
  - ▣ Poor performance with small flows
  - ▣ Really poor performance on wireless networks
  - ▣ Susceptibility to denial of service

# Small Flows

30

- ❑ Problem: TCP is biased against short flows
  - ▣ 1 RTT wasted for connection setup (SYN, SYN/ACK)
  - ▣ *cwnd* always starts at 1
- ❑ Vast majority of Internet traffic is short flows
  - ▣ Mostly HTTP transfers, <100KB
  - ▣ Most TCP flows never leave slow start!
- ❑ Proposed solutions (driven by Google):
  - ▣ Increase initial *cwnd* to 10
  - ▣ TCP Fast Open: use cryptographic hashes to identify receivers, eliminate the need for three-way handshake

# Wireless Networks

31

- ❑ Problem: Tahoe and Reno assume loss = congestion
  - ▣ True on the WAN, bit errors are very rare
  - ▣ False on wireless, interference is very common
- ❑ TCP throughput  $\sim 1/\sqrt{\text{drop rate}}$ 
  - ▣ Even a few interference drops can kill performance
- ❑ Possible solutions:
  - ▣ Break layering, push data link info up to TCP
  - ▣ Use delay-based congestion detection (TCP Vegas)
  - ▣ Explicit congestion notification (ECN)

# Denial of Service

32

- ❑ Problem: TCP connections require state
  - ▣ Initial SYN allocates resources on the server
  - ▣ State must persist for several minutes (RTO)
- ❑ SYN flood: send enough SYNs to a server to allocate all memory/meltdown the kernel
- ❑ Solution: SYN cookies
  - ▣ Idea: don't store initial state on the server
  - ▣ Securely insert state into the SYN/ACK packet (sequence number field)
  - ▣ Client will reflect the state back to the server



DNS



# Layer 8 (The Carbon-based nodes)

34

- ❑ If you want to...
  - ▣ Call someone, you need to ask for their phone number
    - You can't just dial "P R O F G I L L "
  - ▣ Mail someone, you need to get their address first
- ❑ What about the Internet?
  - ▣ If you need to reach Google, you need their IP
  - ▣ Does anyone know Google's IP?
- ❑ Problem:
  - ▣ People can't remember IP addresses
  - ▣ Need human readable names that map to IPs

# Internet Names and Addresses

35

- ❑ Addresses, e.g. 129.10.117.100
  - ❑ Computer usable labels for machines
  - ❑ Conform to structure of the network
- ❑ Names, e.g. [www.northeastern.edu](http://www.northeastern.edu)
  - ❑ Human usable labels for machines
  - ❑ Conform to organizational structure
- ❑ How do you map from one to the other?
  - ❑ Domain Name System (DNS)

# History

36

- ❑ Before DNS, all mappings were in *hosts.txt*
  - ❑ */etc/hosts* on Linux
  - ❑ *C:\Windows\System32\drivers\etc\hosts* on Windows
- ❑ Centralized, manual system
  - ❑ Changes were submitted to SRI via email
  - ❑ Machines periodically FTP new copies of *hosts.txt*
  - ❑ Administrators could pick names at their discretion
  - ❑ Any name was allowed
    - *alans\_server\_at\_sbu\_pwns\_joo\_lol\_kthxbye*

# Towards DNS

37

- ❑ Eventually, the *hosts.txt* system fell apart
  - ❑ Not scalable, SRI couldn't handle the load
  - ❑ Hard to enforce uniqueness of names
    - e.g MIT
      - Massachusetts Institute of Technology?
      - Melbourne Institute of Technology?
  - ❑ Many machines had inaccurate copies of *hosts.txt*
- ❑ Thus, DNS was born

- ❑ DNS Basics
- ❑ DNS Security
- ❑ DNS and Censorship

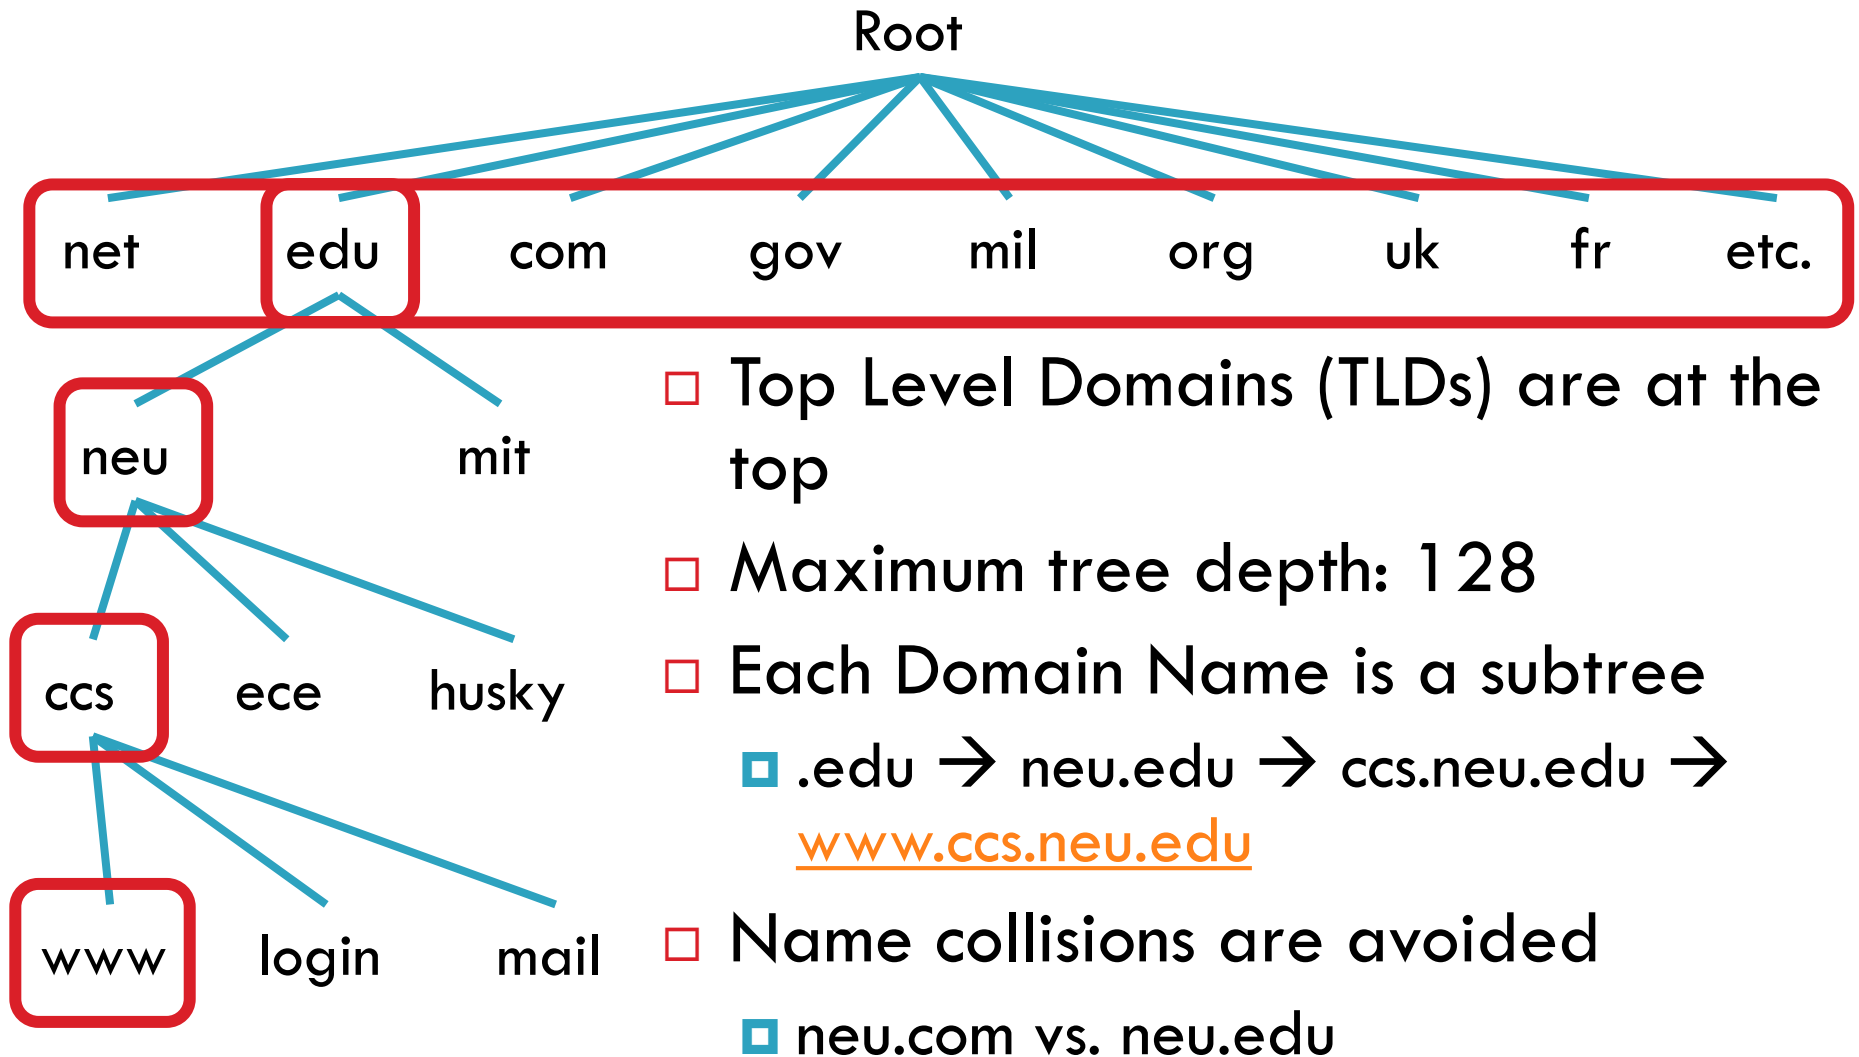
# DNS at a High-Level

39

- ❑ Domain Name System
- ❑ Distributed database
  - ▣ No centralization
- ❑ Simple client/server architecture
  - ▣ UDP port 53, some implementations also use TCP
  - ▣ Why?
- ❑ Hierarchical namespace
  - ▣ As opposed to original, flat namespace
  - ▣ e.g. .com → google.com → mail.google.com

# Naming Hierarchy

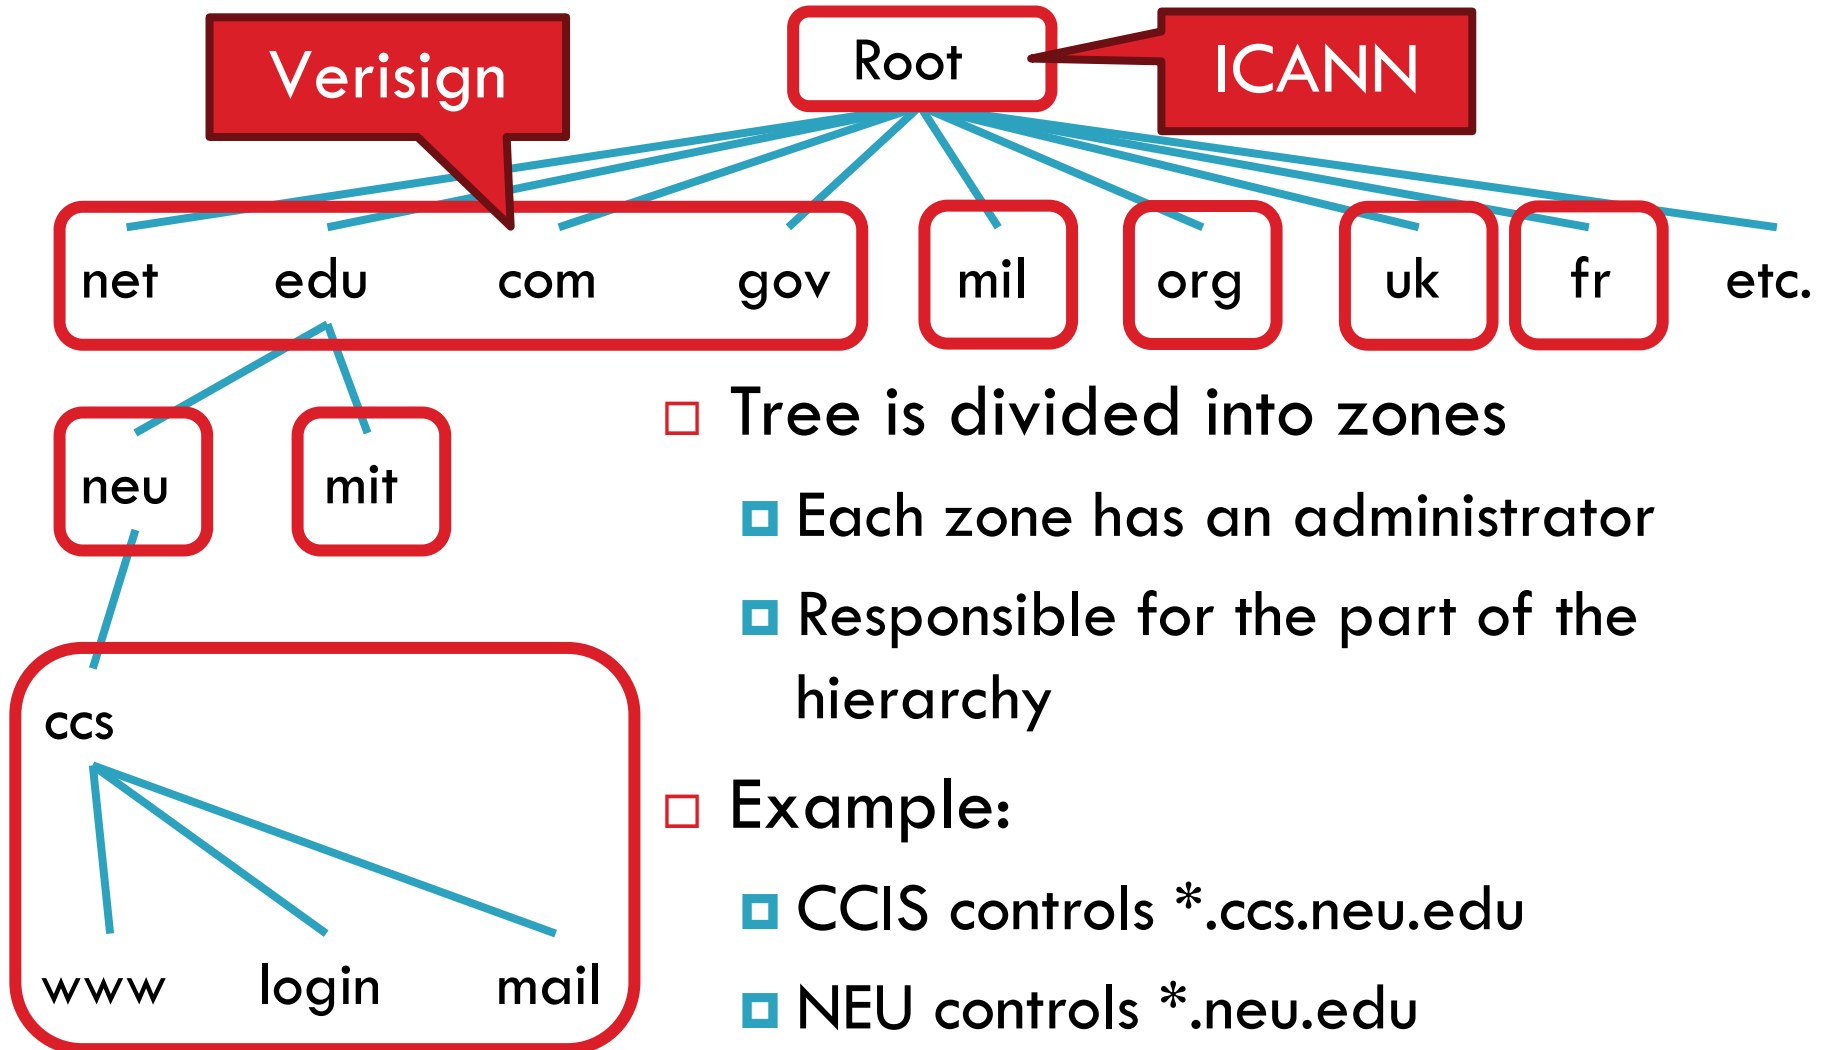
40





# Hierarchical Administration

41



# Server Hierarchy

42

- ❑ Functions of each DNS server:
  - ▣ Authority over a portion of the hierarchy
    - No need to store all DNS names
  - ▣ Store all the records for hosts/domains in its zone
    - May be replicated for robustness
  - ▣ Know the addresses of the root servers
    - Resolve queries for unknown names
- ❑ Root servers know about all TLDs
  - ▣ The buck stops at the root servers

# Root Name Servers

43

- Responsible for the Root Zone File

- ▣ Lists the TLDs and who controls them
- ▣ ~272KB in size

com.	172800	IN	NS	a.gtld-servers.net.
com.	172800	IN	NS	b.gtld-servers.net.
com.	172800	IN	NS	c.gtld-servers.net.

- Administered by ICANN

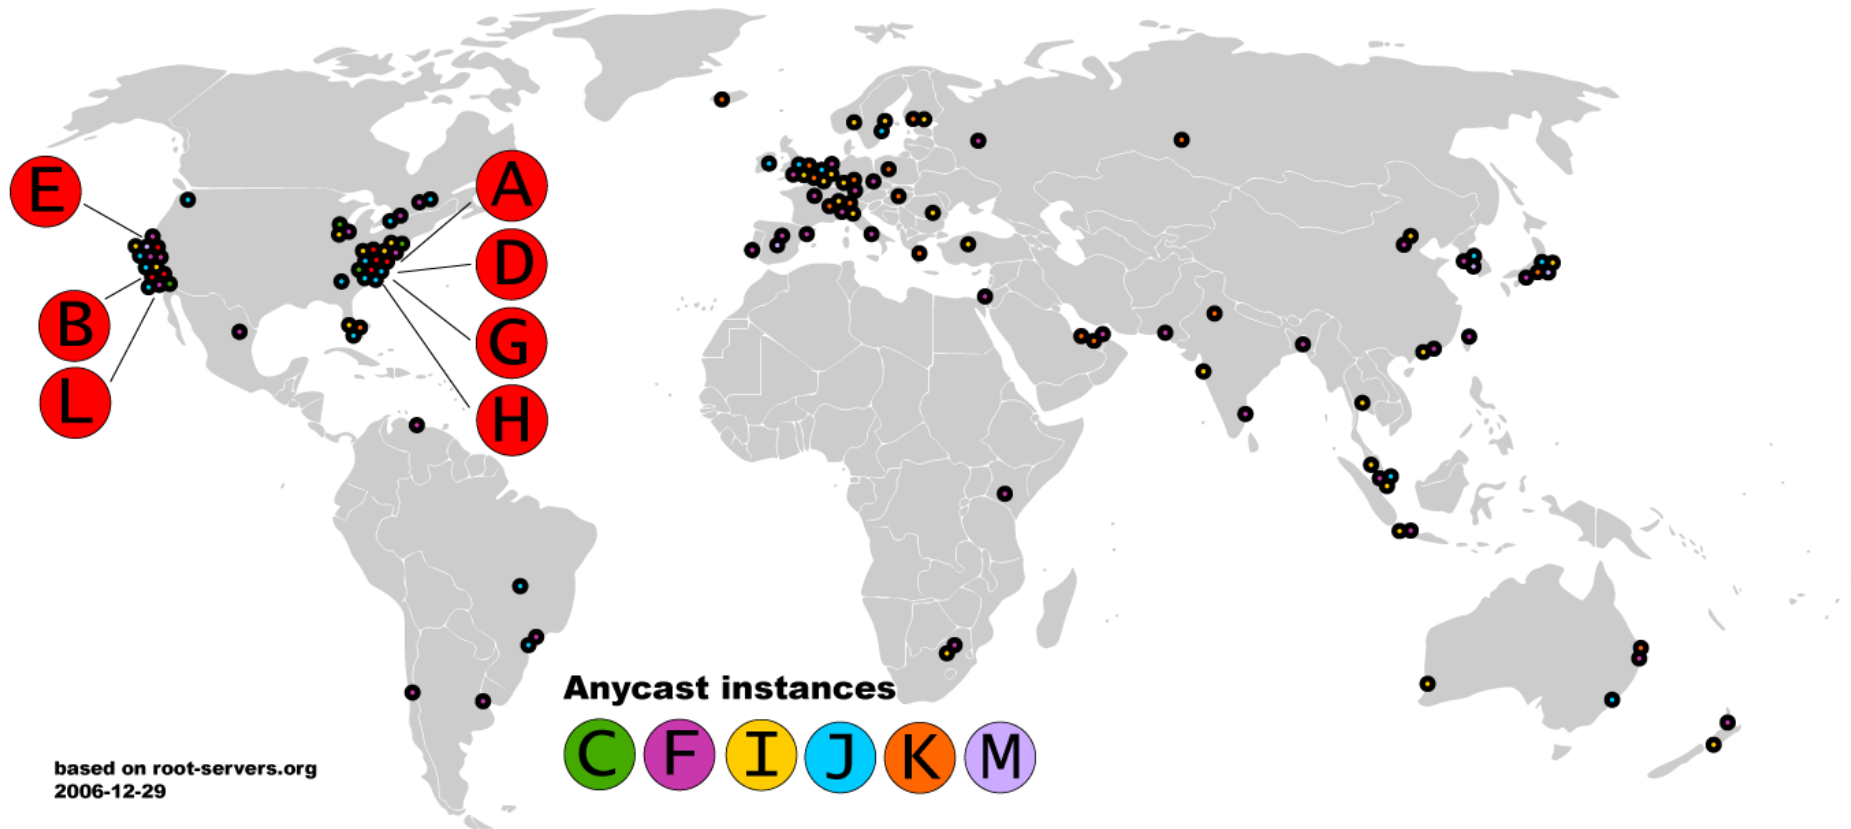
- ▣ 13 root servers, labeled A→M
- ▣ 6 are anycasted, i.e. they are globally replicated

- Contacted when names cannot be resolved

- ▣ In practice, most systems cache this information

# Map of the Roots

44



# Local Name Servers

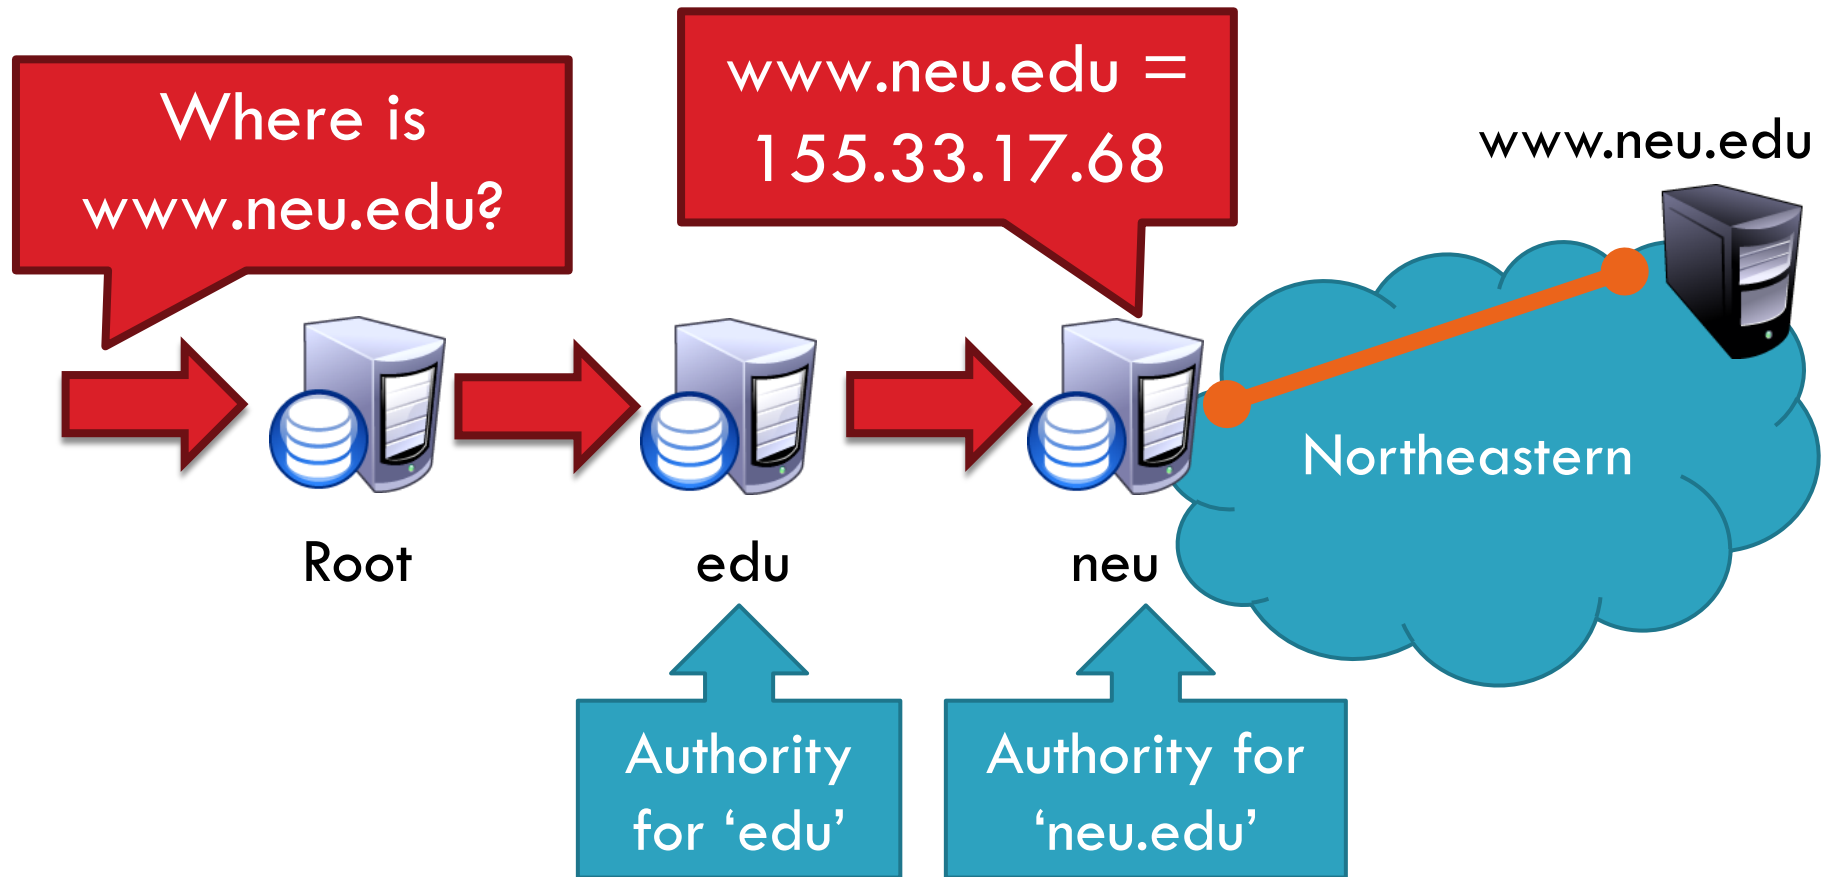
45



- ❑ Each ISP/company has a local, default name server
- ❑ Often configured via DHCP
- ❑ Hosts begin DNS queries by contacting the local name server
- ❑ Frequently cache query results

# Authoritative Name Servers

46



- Stores the name → IP mapping for a given host

# Basic Domain Name Resolution

47


- ❑ Every host knows a local DNS server
  - ▣ Sends all queries to the local DNS server
- ❑ If the local DNS can answer the query, then you're done
  1. Local server is also the authoritative server for that name
  2. Local server has cached the record for that name
- ❑ Otherwise, go down the hierarchy and search for the authoritative name server
  - ▣ Every local DNS server knows the root servers
  - ▣ Use cache to skip steps if possible
    - e.g. skip the root and go directly to .edu if the root file is cached

# Recursive DNS Query

48

Where is [www.google.com](http://www.google.com)?

- ❑ Puts the burden of resolution on the contacted name server
- ❑ How does asgard know who to forward responses too?
  - ▣ Random IDs embedded in DNS queries






# Iterated DNS query

49

Where is [www.google.com](http://www.google.com)?


- Contact server replies with the name of the next authority in the hierarchy
- “I don’t know this name, but this other server might”
- This is how DNS works today



# DNS Propagation

50

- How many of you have purchased a domain name?
  - ▣ Did you notice that it took ~72 hours for your name to become accessible?
  - ▣ This delay is called DNS Propagation



- Why would this process fail for a new DNS name?


# Caching vs. Freshness

51

- DNS Propagation delay is caused by caching



- Zone files may be cached for 1-72 hours



# DNS Resource Records

52

- ❑ DNS queries have two fields: **name** and **type**
- ❑ Resource record is the response to a query
  - ❑ Four fields: (**name**, **value**, **type**, TTL)
  - ❑ There may be multiple records returned for one query
- ❑ What do the **name** and **value** mean?
  - ❑ Depends on the **type** of query and response

# DNS Types

53

## □ Type = A / AAAA

- ▣ Name = domain name
- ▣ Value = IP address
- ▣ A is IPv4, AAAA is IPv6

Query

Name: [www.ccs.neu.edu](http://www.ccs.neu.edu)  
Type: A

Resp.

Name: [www.ccs.neu.edu](http://www.ccs.neu.edu)  
Value: 129.10.116.81

## □ Type = NS

- ▣ Name = partial domain
- ▣ Value = name of DNS server for this domain
- ▣ “Go send your query to this other server”

Query

Name: [ccs.neu.edu](http://ccs.neu.edu)  
Type: NS

Resp.

Name: [ccs.neu.edu](http://ccs.neu.edu)  
Value: 129.10.116.51

# DNS Types, Continued

54

## □ Type = CNAME

- ▣ Name = hostname
- ▣ Value = canonical hostname
- ▣ Useful for aliasing
- ▣ CDNs use this

Query

Name: [foo.mysite.com](http://foo.mysite.com)  
Type: CNAME

Resp.

Name: [foo.mysite.com](http://foo.mysite.com)  
Value: [bar.mysite.com](http://bar.mysite.com)

## □ Type = MX

- ▣ Name = domain in email address
- ▣ Value = canonical name of mail server

Query

Name: [ccs.neu.edu](http://ccs.neu.edu)  
Type: MX

Resp.

Name: [ccs.neu.edu](http://ccs.neu.edu)  
Value: [amber.ccs.neu.edu](http://amber.ccs.neu.edu)

# Reverse Lookups

55

- ❑ What about the IP→name mapping?
- ❑ Separate server hierarchy stores reverse mappings
  - ▣ Rooted at in-addr.arpa and ip6.arpa
- ❑ Additional DNS record **type**: PTR
  - ▣ Name = IP address
  - ▣ Value = domain name
- ❑ Not guaranteed to exist for all IPs

Query

Name: 129.10.116.51  
Type: PTR

Resp.

Name: 129.10.116.51  
Value: [ccs.neu.edu](http://ccs.neu.edu)

# DNS as Indirection Service

56

- ❑ DNS gives us very powerful capabilities
  - ▣ Not only easier for humans to reference machines!
  
- ❑ Changing the IPs of machines becomes trivial
  - ▣ e.g. you want to move your web server to a new host
  - ▣ Just change the DNS record!




# Aliasing and Load Balancing

57

- ❑ One machine can have many aliases

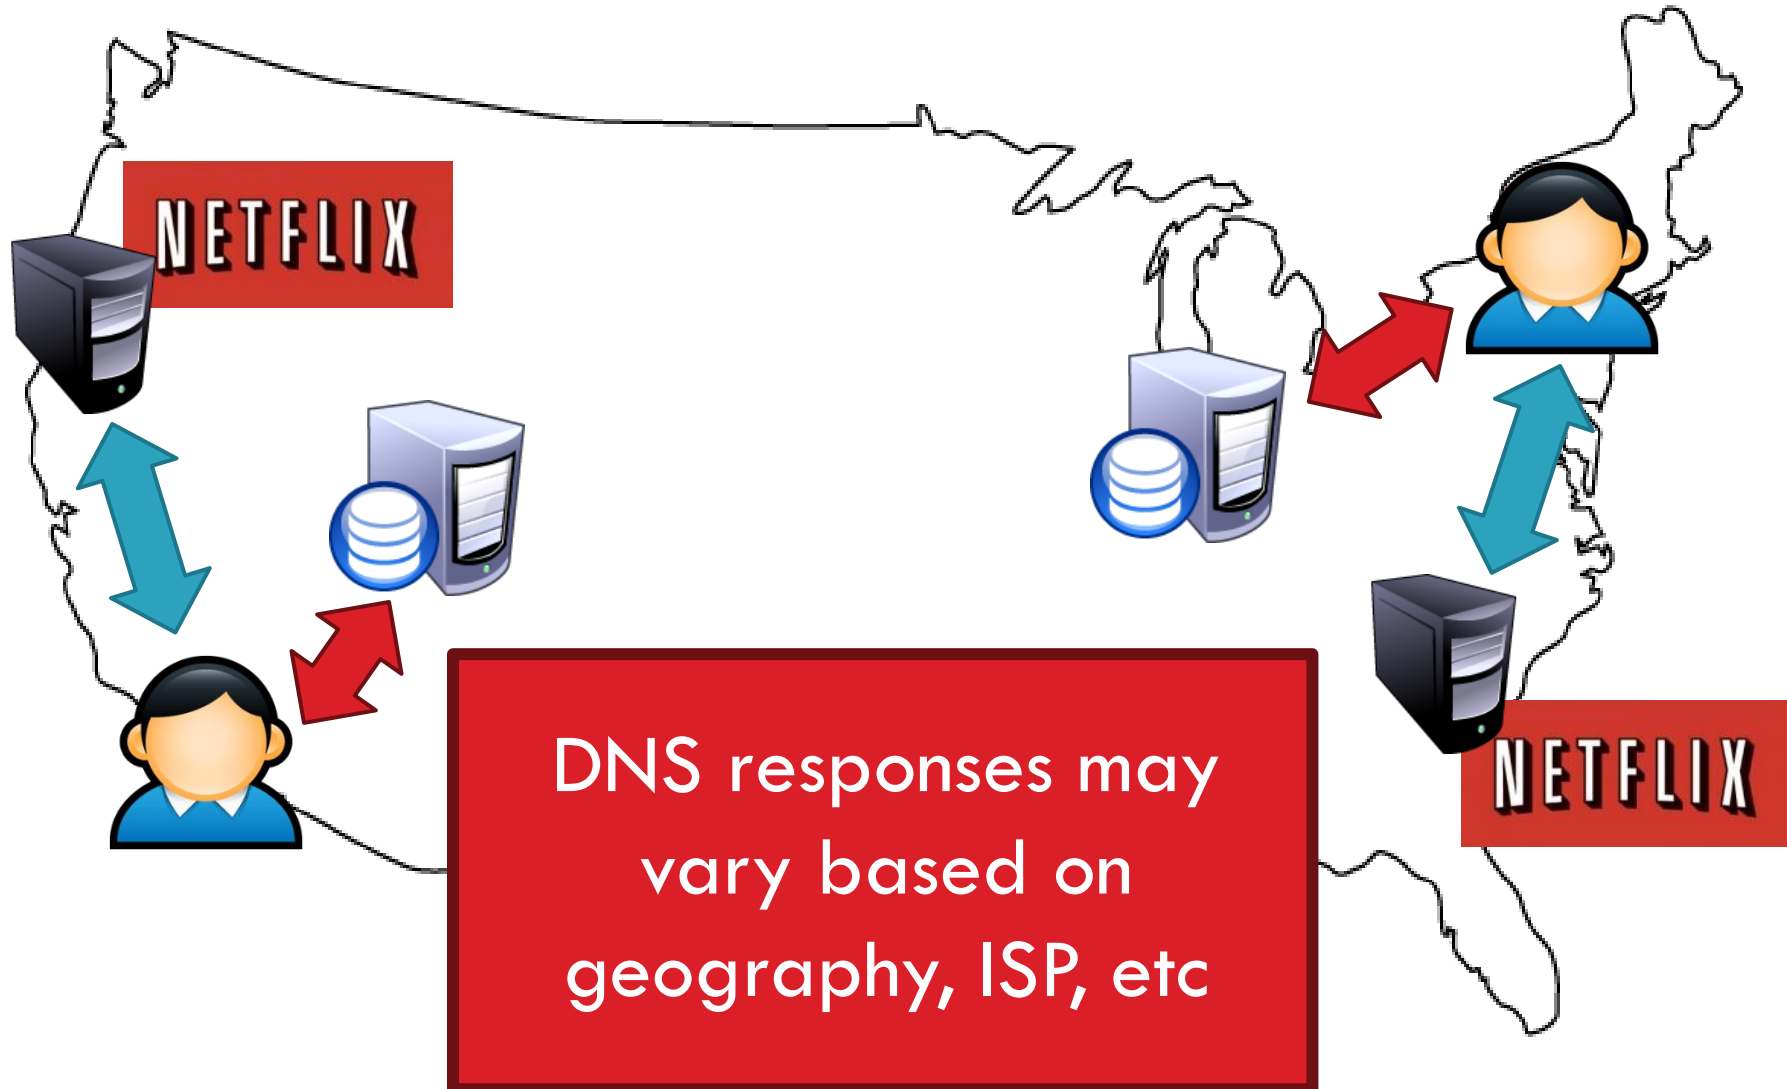


- ❑ One domain can map to multiple machines



# Content Delivery Networks

58



- HTTP Connection Basics
- HTTP Protocol
- Cookies, keeping state + tracking

# Web and HTTP

2-60

*First, a review...*

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`

host name

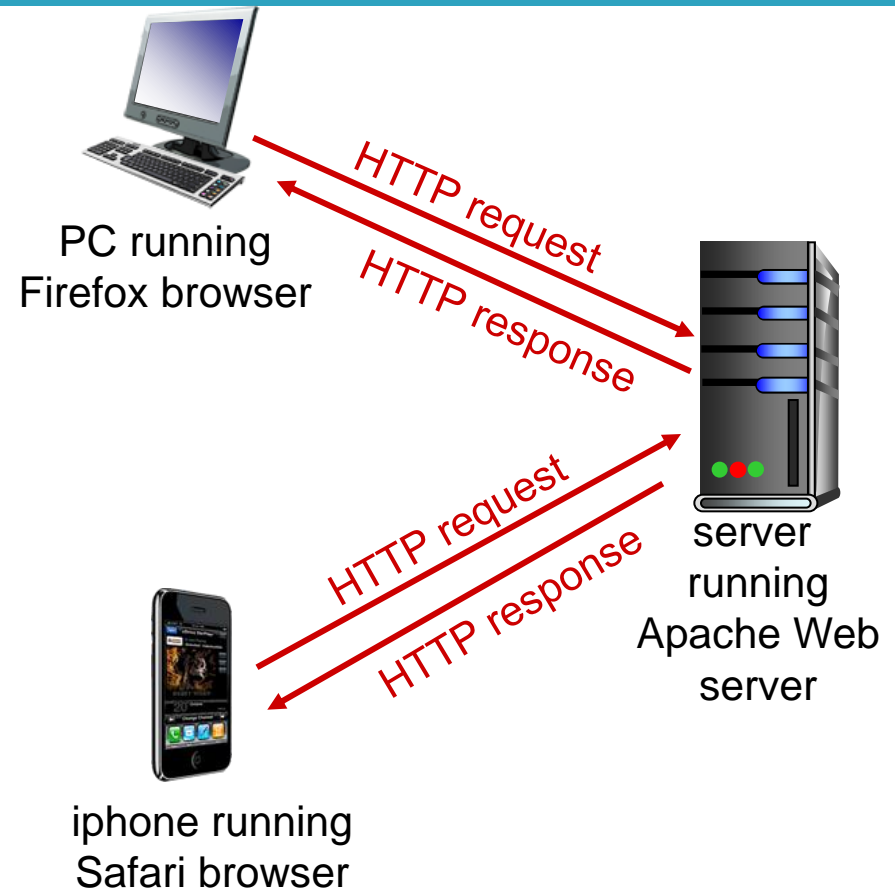
path name

# HTTP overview

2-61

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - ▣ **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
  - ▣ **server**: Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview (continued)

2-62

## *uses TCP:*

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## *HTTP is “stateless” (in theory...)*

- server maintains no information about past client requests

### *aside* protocols that maintain “state” are complex!

- ❖ past history (state) must be maintained
- ❖ if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

2-63

## *non-persistent HTTP*

- at most one object sent over TCP connection
  - ▣ connection then closed
- downloading multiple objects required multiple connections

## *persistent HTTP*

- multiple objects can be sent over single TCP connection between client, server

# Example Web Page

64

page.html

## Harry Potter Movies

As you all know,  
the new HP book  
will be out in June  
and then there will  
be a new movie  
shortly after that...



hpface.jpg


“Harry Potter and  
the Bathtub Ring”



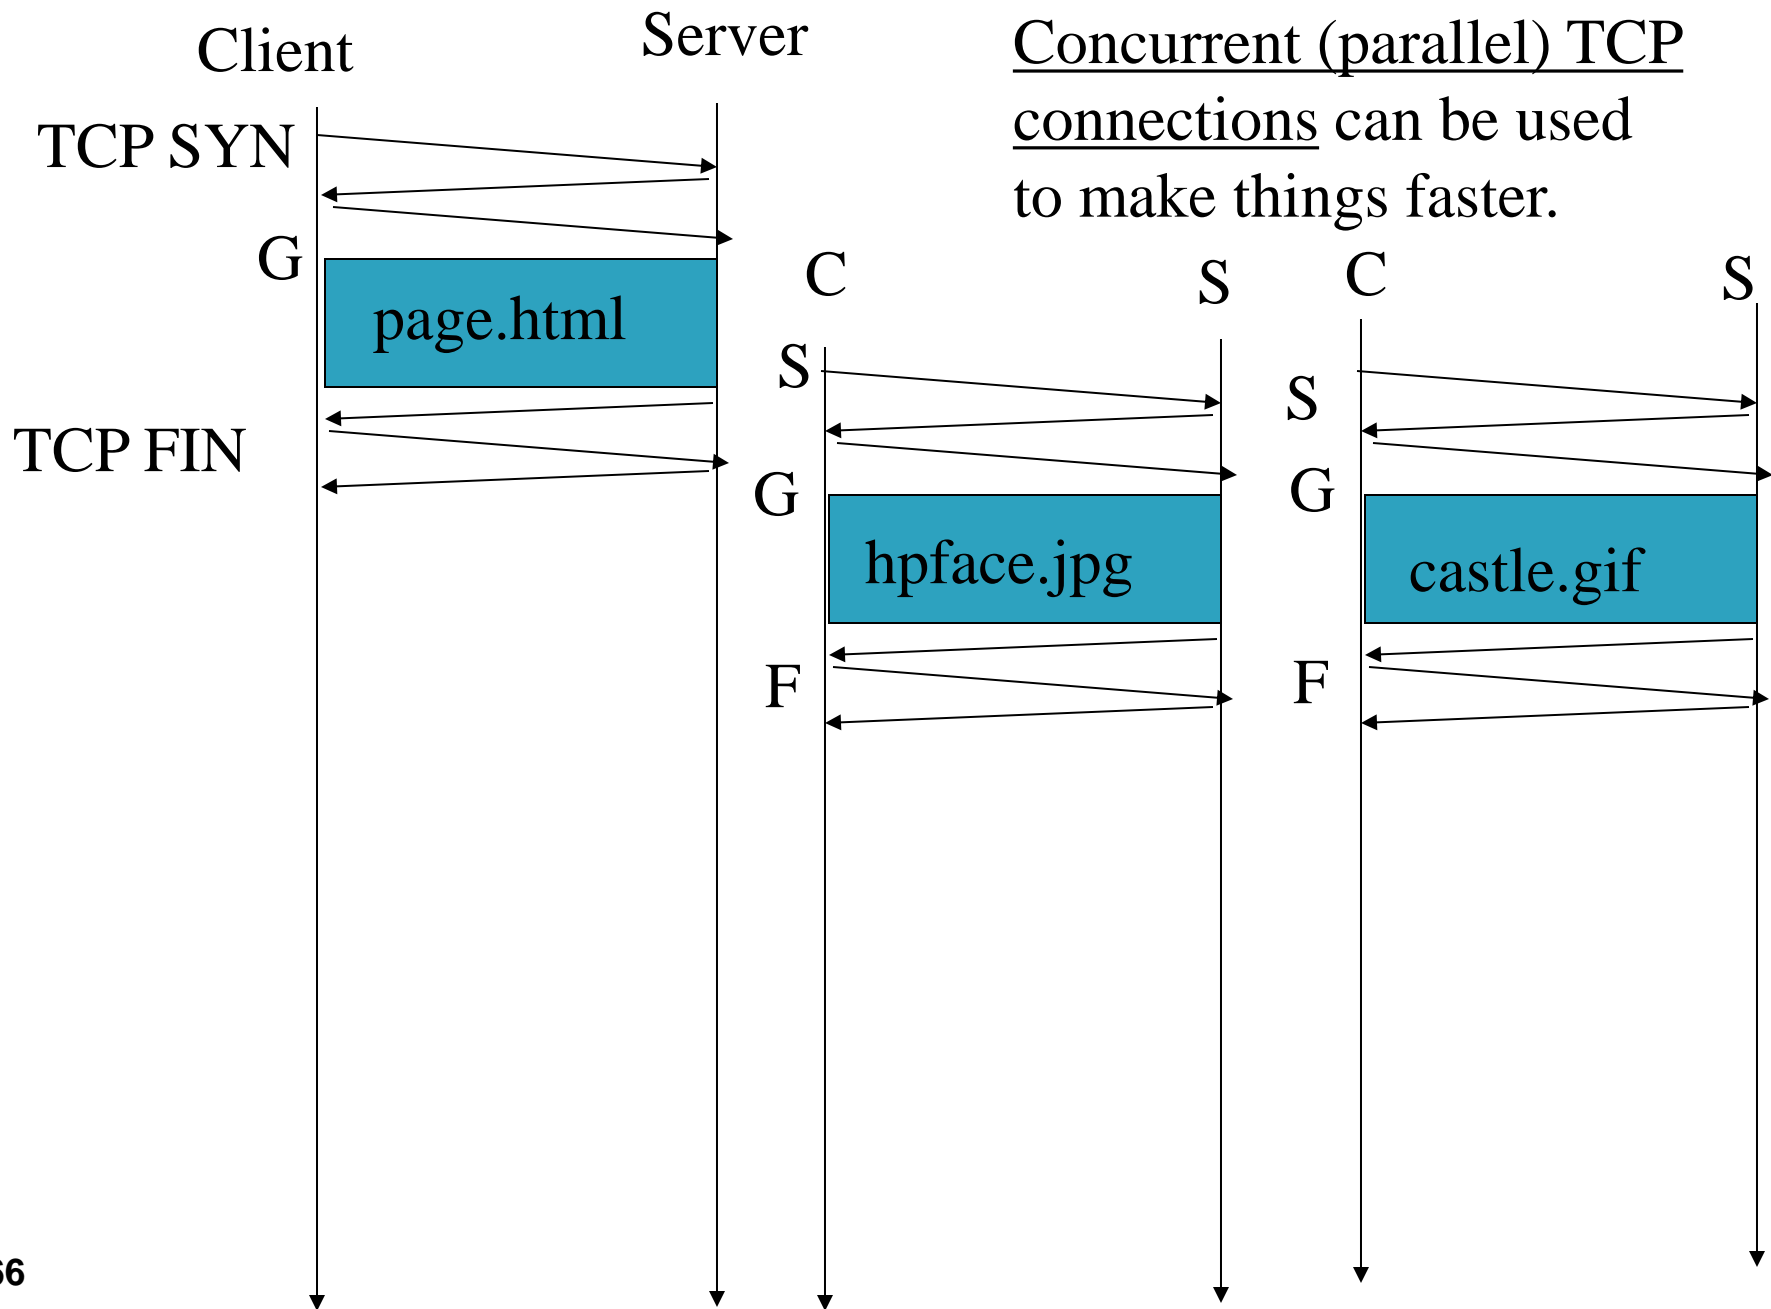
castle.gif



# Non-Persistent HTTP



The “classic” approach in HTTP/1.0 is to use one HTTP request per TCP connection, serially.



# Persistent HTTP

2-67

## *non-persistent HTTP issues:*

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## *persistent HTTP:*

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects


# Non-persistent HTTP: response time

2-68


**RTT:** time for a packet to travel from client to server and back

**HTTP response time:**

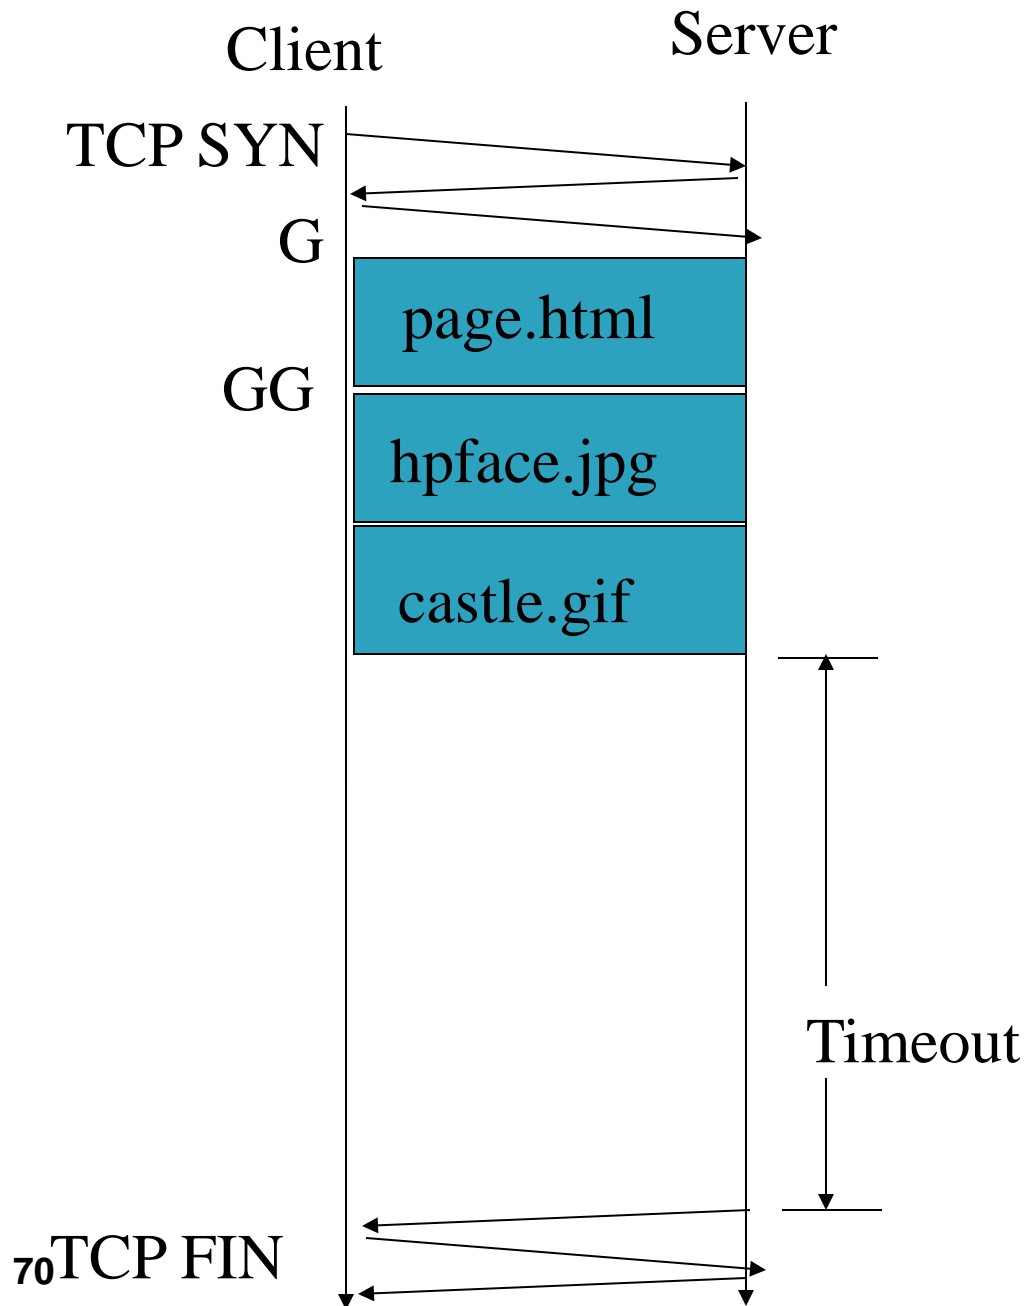
- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
  - ▣ This assumes HTTP GET piggy backed on the ACK
- file transmission time
- non-persistent HTTP response time =  
 $2\text{RTT} + \text{file transmission time}$



# Persistent HTTP



The “persistent HTTP” approach can re-use the same TCP connection for Multiple HTTP transfers, one after another, serially. Amortizes TCP overhead, but maintains TCP state longer at server.



The “pipelining” feature in HTTP/1.1 allows requests to be issued asynchronously on a persistent connection. Requests must be processed in proper order. Can do clever packaging.

- HTTP Connection Basics
- HTTP Protocol
- Cookies, keeping state + tracking

# HTTP request message

2-72

- two types of HTTP messages: *request, response*
- **HTTP request message:**
  - ▣ ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

header  
lines

carriage return,  
line feed at start  
of line indicates  
end of header lines  
Application Layer

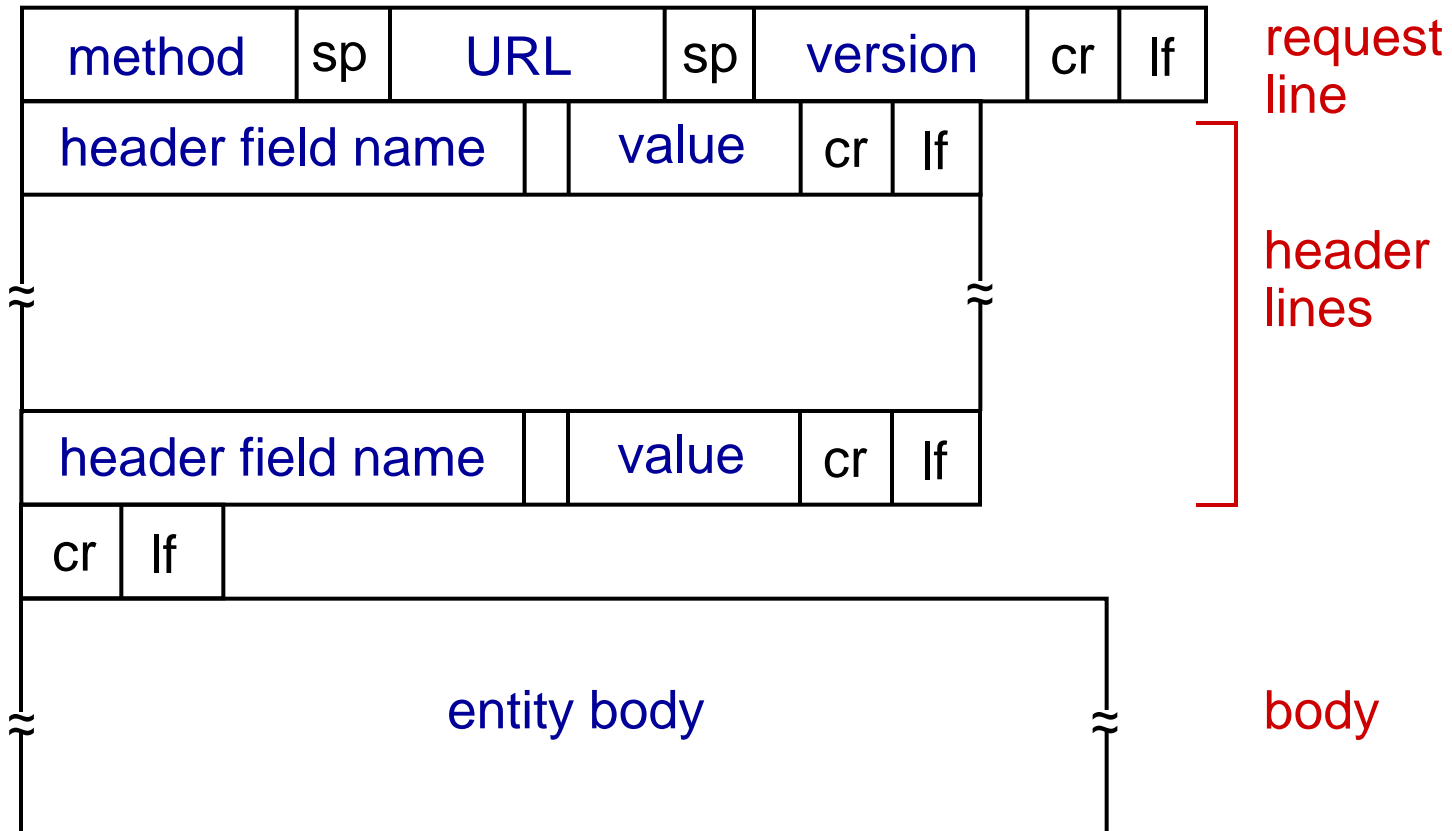
carriage return character  
line-feed character

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```



# HTTP request message: general format

2-73



# Uploading form input

2-74

## POST method:

- ❑ web page often includes form input
- ❑ input is uploaded to server in entity body

## URL method:

- ❑ uses GET method
- ❑ input is uploaded in URL field of request line:  
`www.somesite.com/animalsearch?monkeys&banana`

# Method types

2-75

## HTTP/1.0:

- GET
- POST
- HEAD
  - ▣ asks server to leave requested object out of response

## HTTP/1.1:

- GET, POST, HEAD
- PUT
  - ▣ uploads file in entity body to path specified in URL field
- DELETE
  - ▣ deletes file specified in the URL field

# HTTP response message

2-76

status line

(protocol

status code

status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
      GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
      1\r\n
\r\n
data data data data data ...
```

Application Layer

# HTTP response status codes

2-77

- ❖ status code appears in 1st line in server-to-client response message.

- ❖ some sample codes:

**200 OK**

- ▣ request succeeded, requested object later in this msg

**301 Moved Permanently**

- ▣ requested object moved, new location specified later in this msg  
(Location:)

**400 Bad Request**

- ▣ request msg not understood by server

**404 Not Found**

- ▣ requested document not found on this server

**505 HTTP Version Not Supported**

# Trying out HTTP (client side) for yourself

2-78

1. Telnet to your favorite Web server:

```
telnet cis.poly.edu 80
```

opens TCP connection to port 80  
(default HTTP server port) at cis.poly.edu.  
anything typed in sent  
to port 80 at cis.poly.edu

2. type in a GET HTTP request:

```
GET /~ross/ HTTP/1.1  
Host: cis.poly.edu
```

by typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

3. look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)

- HTTP Connection Basics
- HTTP Protocol
- Cookies, keeping state + tracking

# User-server state: cookies

2-80

many Web sites use cookies

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

**example:**

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - ▣ unique ID
  - ▣ entry in backend database for ID



# Cookies: keeping “state” (cont.)

2-81

client



server



cookie file

usual http request msg

Amazon server  
creates ID  
1678 for user

usual http response  
**set-cookie: 1678**

create  
entry

backend  
database




ebay 8734  
amazon 1678

usual http request msg  
**cookie: 1678**

cookie-  
specific  
action

access

usual http response msg



access

cookie-  
specific  
action

one week later:



ebay 8734  
amazon 1678

usual http request msg  
**cookie: 1678**

usual http response msg

Application Layer

# Cookies (continued)

2-82

*what cookies can be used for:*

- ☐ authorization
- ☐ shopping carts
- ☐ recommendations
- ☐ user session state (Web e-mail)

*cookies and privacy:*

- ❖ cookies permit sites to learn a lot about you
- ❖ you may supply name and e-mail to sites

aside

*how to keep “state”:*


- ❖ protocol endpoints: maintain state at sender/receiver over multiple transactions
- ❖ cookies: http messages carry state

# Cookies + Third Parties

83

- Example page (from Wired.com)

## Elijah Wood's New Movie Is a Prophetic Thriller About Celebrity Hacking



Elijah Wood in *Open Windows*.  courtesy Cinedigm

# How it works

84

And it's not just Facebook!



Wi

GET article.html



GET sharebutton.gif  
Cookie: FBCOOKIE



Facebook now knows you visited this Wired article.  
Works for all pages where 'like'/'share' button is embedded!